

CS542200 Parallel Programming

Homework 5: Observe the behavior of UCX yourself

Due: Sun. Dec 21 23:59, 2025

1 GOAL

Develop observability code within the UCX framework to gain a comprehensive understanding of its architecture and enhance our knowledge in modern parallel computing frameworks, with the following objectives:

1. **Comprehend the UCX Framework:** Delve into the UCX framework to understand the roles and interconnections of its subcomponents, thereby gaining a deeper insight into its overall functionality.
2. **Examine TLS Impact:** Investigate how Transport Layer Security (TLS) is influenced by the interactions and dependencies between UCP, UCS, and UCT within the UCX framework.
3. **Learn Program Structure:** Acquire knowledge about the program structure of modern parallel computing frameworks, focusing on how they are designed, implemented, and optimized for performance.

In this HW, students are advised to copy the [Report Template](#) first and complete the assignment step by step, following the question in the report.

2 REQUIREMENTS

-
- In this assignment, you are required to modify and recompile the UCX system (the underlying communication protocol for MPI) to add the following two features:
 - Display the UCX transport protocols currently configured in the system.
 - ◆ You are required to insert code at // TODO: PP-HW (src/ucs/config/parser.c) to ensure your program is compatible with the UCX framework.
 - Output the UCX transport protocols actively utilized by the program.

3 RUN YOUR PROGRESS

-
1. Recompile UCX and install it in the `$HOME/ucx-pp` directory, ensuring that a `lib` folder is present that contains the required dynamic link files. (`libucp.so.0`,

`libuct.so.0, libucm.so.0, libucs.so.0).`

Note that every time you modify the UCX code, you need to re-execute the last line of commands to recompile.

```
git clone -b pp2024 https://github.com/NTHU-LSALAB/UCX-lsalab
cd UCX-lsalab
mkdir build && cd build
../configure --prefix=$HOME/ucx-pp/ --with-go=no
srun -n 1 -c 12 make -j12 install
```

2. We provide a custom command `mpiucx`, designed to replace `mpirun`, which will utilize your self-compiled UCX as the underlying communication framework for running MPI parallel programs. You can find the test cases in `UCX-lsalab/test/`

```
module load openmpi/ucx-pp
mpiucx -x UCX_LOG_LEVEL=info -np 2 ./mpi_hello.out
```

3. Alter the internal code of UCX to enable printing of:
 - The UCX_TLS information is currently specified by the system.
 - The final TLS transport method was selected by UCX.
4. You may need to trace the files within `src/ucp`, `src/uct`, and `src/ucs` directories and modify at least the following two files to complete the homework.
 - `ucp/core/ucp_worker.c`
 - Invoke `ucp_config_print` to print `UCX_TLS`.
 - Print Line 2 by identifying a key variable.
 - `ucs/config/parser.c`
 - Implement the `ucs_config_parser_print_opts` function to be invoked by `ucp_config_print` here.

4 RESULT

For each transport, the output comprises 2 lines.

- Line 1: The information from `UCX_TLS` needs to be exactly the same.
- Line 2: The information should include only the key strings of transport protocols selected by UCX. The example must contain strings.
`cfg#0 tag(sysv/memory cma/memory)`

Sample input & output

- For the openmpi/ucx-pp configuration, the default transport protocol is set to utilize ud_verbs.

```
sky@apollo01:~/UCX-lsalab/test/mpi$ mpiucx -np 1 ./mpi_hello.out
UCX_TLS=ud_verbs
0x55c7c5a66f70 self cfg#0 tag(ud_verbs(ibp3s0:1))
Hello world from processor apollo01, rank 0 out of 1 processors
```

- We have enabled UCX to leverage all available transport protocols by setting up UCX_TLS.

```
sky@apollo01:~/UCX-lsalab/test/mpi$ mpiucx -n 2 -x UCX_TLS=all ./send_recv.out
[apollo01:4170706] SET UCX_TLS=all
UCX_TLS=all
0x562418e25fb0 self cfg#0 tag(self/memory cma/memory)
UCX_TLS=all
0x55e8a8efa070 self cfg#0 tag(self/memory cma/memory)
UCX_TLS=all
0x562418e25fb0 intra-node cfg#1 tag(sysv/memory cma/memory)
UCX_TLS=all
0x55e8a8efa070 intra-node cfg#1 tag(sysv/memory cma/memory)
Process 1 received message 'Hello from rank 0' from process 0
Process 0 sent message 'Hello from rank 0' to process 1
```

5 REPORT

Here is the report template <https://hackmd.io/@khWVI90iQ7iha6nHAgInUA/SyppEw-l-e>
 Please download the markdown and submit the .md file.

6 GRADING

-
- [30%] Correctness**
 - The pass rate of test data for ucx-judge.
 - [20%] Demo**
 - A demo session will be held remotely. You'll be asked questions about the homework.
 - [50%] Report**

- Grading is based on your evaluation, discussion and writing. If you want to get more points, design or conduct more experiments to analyze your implementation.

7 SUBMISSION

Upload the files below to eeclass. (**DO NOT COMPRESS THEM**)

- hw5.diff
 - You can obtain this file by
 - `git add -A && git commit -m "Observed the behavior of UCX."`
 - `git diff --color remotes/origin/pp2024 > hw5.diff`
- hw5_{student_ID}.md or hw5_{student_ID}.pdf
 - If possible, please submit it in Markdown and provide a HackMD link or PDF file as a backup.
 - Please confirm that your image is displayed correctly.
- **KEEP `$HOME/ucx-pp/` CORRECTLY**

8 HINTS

- Type `ucx-judge` on **Origo Apollo** to run the test cases.
- Scoreboard:
 - <https://pp.lsab.cs.nthu.edu.tw/ucx/>
- You can use:
 - `ucx_info -d` to show all available TLS in your cluster.
 - `-x UCX_LOG_LEVEL=info` in `mpirun` to obtain debug information during UCX runtime
 - `-x UCX_TLS=<TLS>` to specify the permitted TLS for your program.
- Refer to slide **p37** of the presentation to determine which component is accountable for the `UCX_TLS` configuration in the system.
- Consult slides **p41** to **p42** of the presentation for insights into UCP's architecture, which will aid in navigating your code to pinpoint the transport protocols actually in use.

- If you encounter any problems with the homework specification, judge scripts, example source code, or test cases, please contact the TA at pp@lsalab.cs.nthu.edu.tw or eeclasse.
- You are allowed to discuss and exchange ideas with others, but you are required to write the code yourself. If we find you cheating, you'll receive 0 points.