


# AMD GPU Introduction

PP 2025 LAB5



國立清華大學叢集電腦競賽團隊  
Student Cluster Competition Team of NTHU



# Outline

Basic Knowledge of AMD GPU

ROCM API

ROCM Profiler



# Basic Knowledge of AMD GPU



# What is Thread, Block, Grid?

## Thread

The basic unit of work executed by the GPU.

Threads are group in **warp**, threads in the same warp must execute the same instruction simultaneously.



## Block

Block(thread block) is a group of threads that can cooperate and communicate with each other.


All threads in a single block are mapped to and executed on the same **CU**(compute unit).



## Grid

The highest-level organizational structure for threads executing a kernel function, representing a single launch of that kernel.

A grid consists of multiple thread blocks, which execute the kernel concurrently.




# What is Warp and Cu?

- Warp
  - 32/64 threads per warp in amd gpu
  - threads in same warp execute the same instruction simultaneously
  - avoid warp divergence is important
- Cu(Compute Unit)
  - each Cu can handle multiple thread blocks
  - but holds only limited number of warps
  - Cu shares the same space of L1 cache and shared memory



# Memory Hierarchy



## Thread Wide

This is the lowest and most confined level of the hierarchy.

- Data is private to a single thread and is stored in registers.
- A thread cannot see or access the registers of other threads.
- Actions of one thread have no effect or visibility to any other thread at this level.

## Block Wide

A block is a group of threads that execute together on the same CU.

- Threads in the block can cooperate and share data.
- Achieved through shared memory and synchronization.

## Device Wide

This is the highest level of the hierarchy and includes all threads and all blocks running on the entire GPU.

- All threads and blocks can access global memory.
- Data sharing through L2 cache and global memory.



# Some Different with Cuda

- NVIDIA GPUs have a fixed warp size of 32 threads.
- AMD GPUs typically use a warp size of 64 threads, though some architectures are 32.
- NVIDIA allows the shared-memory size per SM to be configured dynamically, whereas AMD's is fixed at 64 KB per Cu.



# ROCM API



國立清華大學叢集電腦競賽團隊  
Student Cluster Competition Team of NTHU

# HIP Compiler and Header

Hip compiler: `hipcc`

Hip runtime header: `include <hip/hip_runtime.h>`

Compile hip code: `hipcc -o <xxx> <xxx.hip>`



# SMI

## rocm-smi

ROCM System Management Interface													
Concise Info													
Device	Node	IDs (DID, GUID)	Temp (Edge)	Power (Avg)	Partitions (Mem, Compute, ID)	SCLK	MCLK	Fan	Perf	PwrCap	VRAM%	GPU%	
0	11	0x738c, 23480	32.0°C	34.0W	N/A, N/A, 0	300Mhz	1200Mhz	20.78%	auto	290.0W	0%	0%	
1	10	0x738c, 64802	31.0°C	34.0W	N/A, N/A, 0	300Mhz	1200Mhz	20.78%	auto	290.0W	0%	0%	
2	9	0x738c, 59802	31.0°C	34.0W	N/A, N/A, 0	300Mhz	1200Mhz	20.78%	auto	290.0W	0%	0%	
3	8	0x738c, 44650	31.0°C	34.0W	N/A, N/A, 0	300Mhz	1200Mhz	20.78%	auto	290.0W	1%	0%	
4	13	0x738c, 30589	33.0°C	39.0W	N/A, N/A, 0	300Mhz	1200Mhz	20.78%	auto	290.0W	0%	0%	
5	12	0x738c, 15436	31.0°C	36.0W	N/A, N/A, 0	300Mhz	1200Mhz	20.78%	auto	290.0W	0%	0%	

End of ROCm SMI Log



# GPU Info

## rocminfo

```
Workgroup Max Size:      1024(0x400)
Workgroup Max Size per Dimension:
  x                      1024(0x400)
  y                      1024(0x400)
  z                      1024(0x400)
Grid Max Size:           4294967295(0xffffffff)
Grid Max Size per Dimension:
  x                      4294967295(0xffffffff)
  y                      4294967295(0xffffffff)
  z                      4294967295(0xffffffff)
FBarrier Max Size:       32
```



# Allocate Device Global Memory

- **hipError\_t hipMalloc(void\*\* devPtr, size\_t size);**
  - **devPtr** – pointer to a pointer that will receive the device address.
  - **size** – number of bytes to allocate in GPU global memory.
  - Returns a **hipError\_t** status code.
- **hipError\_t hipMemcpy**  
**(void\* dst, const void\* src, size\_t size, hipMemcpyKind kind);**
  - **dst / src** – destination and source pointers.
  - **size** – number of bytes to copy.
  - **kind** – direction of the copy.
  - Returns a **hipError\_t** status code.

Don't forget to **hipFree!**



# Allocate Device Global Memory

Device Pointer

Host Pointer

```
void solve(const float* logits, const int* true_labels, float* loss, int N, int C) {
```

```
    float *d_logits, *d_loss;  
    int *d_true_labels;
```

```
    hipMalloc(&d_logits, N * C * sizeof(float));  
    hipMalloc(&d_true_labels, N * sizeof(int));  
    hipMalloc(&d_loss, sizeof(float));
```

```
    hipMemcpy(d_logits, logits, N * C * sizeof(float), hipMemcpyHostToDevice);  
    hipMemcpy(d_true_labels, true_labels, N * sizeof(int), hipMemcpyHostToDevice);  
    hipMemcpy(d_loss, loss, sizeof(float), hipMemcpyHostToDevice);
```

Memory copy from host to device

```
    hipMemcpy(loss, d_loss, sizeof(float), hipMemcpyDeviceToHost);
```

Memory copy from device to host



# Launch Kernel Function

```
hipError_t hipLaunchKernelGGL(  
    kernelFunc,           // __global__ kernel name  
    dim3 gridDim,         // grid size (x,y,z)  
    dim3 blockDim,        // block size (x,y,z)  
    size_t sharedMemBytes, // optional dynamic shared memory (bytes)  
    hipStream_t stream,   // stream handle (0 for default)  
    kernelArgs...          // arguments to the kernel  
);
```

**OR**

necessary	optional
kernelName<<<gridDim, blockDim,	sharedMemBytes, stream>>>(arg1, arg2, ...);



# Other APIs

Most HIP APIs work the same as CUDA—just press Ctrl + F and replace `cuda` with `hip`.



# ROCM Profiler



國立清華大學叢集電腦競賽團隊  
Student Cluster Competition Team of NTHU

# ROCProf Compute

rocprof-compute **profile** -n <xxx> -- ./<xxx> <args>

rocprof-compute **analyze** -p workloads/<xxx>/MI100/

- profile - run and record profiling data
  - --name (-n) : The profiler will create a folder like workloads/<xxx>/
  - --kernel (-k) : profile specific kernel
- analyze - open and inspect data from a previous run
  - --path (-p) : inspect data generate by profile
- database - import/export data for Grafana or other database-backed viewers

<https://rocm.docs.amd.com/projects/rocprofiler-compute/en/latest/how-to/use.html>



# ROCProf Compute

## 0.1 Top Kernels

	Kernel_Name	Count	Sum(ns)	Mean(ns)	Median(ns)	Pct
0	phase3(int, int*, int) [clone .kd]	157.00	422091878.00	2688483.30	2688487.00	96.16
1	phase2(int, int*, int) [clone .kd]	157.00	11097161.00	70682.55	71200.00	2.53
2	phase1(int, int*, int) [clone .kd]	157.00	5131208.00	32682.85	29920.00	1.17
3	initializeDeviceDist(int*, int) [clone .kd]	1.00	601441.00	601441.00	601441.00	0.14
4	updateDeviceDist(int*, int const*, int, int) [clone .kd]	1.00	20480.00	20480.00	20480.00	0.00



# ROCProf Compute

gpu_model	MI100
gpu_arch	gfx908
gpu_l1	16
gpu_l2	8192
cu_per_gpu	120
simd_per_cu	4
se_per_gpu	8
wave_size	64
workgroup_max_size	1024

Size of L1 Cache(KB) per CU

Size of L2 Cache(KB) in whole GPU

Number of Cus in whole GPU

wave = warp

workgroup = threadblock



# ROCProf Compute

```
rocprof-compute analyze -p workloads/<baseline>/MI100/ \
-p workloads/<compare>/MI100/
```

0. Top Stats						
0.1 Top Kernels						
	Kernel_Name	Count	Count	Abs Diff	Sum(ns)	Sum(ns)
0	phase3(int, int*, int) [clone .kd]	157.00	235.0 (49.68%)	78.00	421864145.00	941825411.0 (123.25%)
1	phase2(int, int*, int) [clone .kd]	157.00	235.0 (49.68%)	78.00	11090903.00	17447075.0 (57.31%)
2	phase1(int, int*, int) [clone .kd]	157.00	235.0 (49.68%)	78.00	5124330.00	7125766.0 (39.06%)
3	initializeDeviceDist(int*, int) [clone .kd]	1.00	1.0 (0.0%)	0.00	603361.00	1368001.0 (126.73%)
4	updateDeviceDist(int*, int const*, int, int) [clone .kd]	1.00	1.0 (0.0%)	0.00	20480.00	8000.0 (-60.94%)



# ROCProfv2

```
rocprofv2 -i <xxx(metric.txt)> -o <xxx> -d <xxx> ./<xxx> <args>
```


- --input (-i) : specify a input file for metric profiling
- metric.txt : your metric input file
  - Use **rocprofv2 --list-counters** to check all the usable metrics
  - pmc : The rows in the text file beginning with **pmc:** are the group of metrics you are interested in collecting
  - Example metric file: cat > metrics.txt <<'EOF' pmc: SQ\_WAVES EOF
- --output-file-name (-o) : profile output csv filename
- --output-directory (-d) : profile output directory
- You will get a result\_<xxx>.csv, download it and check the profile result



# ROCm Systems Profiler

rocprof-sys-run --trace -- ./<xxx> <args>

Download this : [perfetto-trace-1256288.proto](#)

Open it on [perfetto](#) : 

You will get :



# SPEC



國立清華大學叢集電腦競賽團隊  
Student Cluster Competition Team of NTHU



# THE END

